

LLM, RAG – Vector Stores

Qdrant:

qdrant stands out in the vector database landscape with its unique set of features:

Advanced indexing mechanisms for optimized search operations.

Scalability to manage large datasets without compromising on speed.

qdrant showcases exceptional performance when dealing with massive volumes of high-dimensional data. Its efficient indexing strategies make it a top choice for applications demanding real-time search capabilities.

pgvector:

pgvector is a versatile vector database known for its robust features and performance capabilities. Its key features include:

Efficient storage of high-dimensional vectors.

Seamless integration with PostgreSQL, enhancing existing database functionalities.

Where pgvector truly shines is in its ability to handle complex data structures with ease. It excels in scenarios requiring quick retrieval of similar vectors, making it a valuable asset for applications like recommendation systems and image recognition.

FAISS:

FAISS (Facebook AI Similarity Search) is a library for efficient similarity search and clustering of dense vectors.

It can search multimedia documents (e.g. images) in ways that are inefficient or impossible with standard database engines (SQL).

It contains algorithms that search in sets of vectors of any size, up to ones that possibly do not fit in RAM. It also contains supporting code for evaluation and parameter tuning.